



Four-class episode distributions across AdSERP [LAB]

clicked · deferred · evaluated-rejected — per-trial, per-rank, within-trial, and co-occurrence


(a) Per-trial episode counts — how thick is each class per session?




(b) Class prevalence by rank — where does each class live on the page?



(c) Episode position within trial — early, middle, or late in the sequence?



(d) Class co-occurrence within trial — homogeneous or mixed?



Classes via NB22 regression rule: approached AND NOT clicked AND gaze_regression_label (eye-tracked; scroll telemetry not used for labels). N = 2,287 trials, 4,583 episodes.